YUXUAN WANG

Contact: ethan.yuxuan.wang@gmail.com · Homepage: www.yuxuanw.me ·

EDUCATION

Nanyang Technological University (NTU), Computer Science and Engineering,
Ph.D. Candidate — Aug. 2022 - Present

Supervised by Prof. *Hanwang Zhang* in Computer Vision.
Research Topic: *3D Generaton*, *3D Controllable Editing*, and *Vision-Language Understanding*

National University of Singapore (NUS), Electrical and Computer Engineering,
Master of Science (*GPA: 4.47/5*) — Aug. 2021 - Jun. 2022

Supervised by Prof. *Mike Zheng Shou* in Computer Vision.

Beihang University, Electronic Information Engineering,
Bachelor of Engineering (*GPA: 87.5/100*) — Sep. 2016 - Jun. 2020

PUBLICATIONS

Personalize Your Gaussian: Consistent 3D Scene Personalization from a Single Image
Yuxuan Wang, Xuanyu Yi, Qingshan Xu, Yuan Zhou, Long Chen, Hanwang Zhang
Arxiv Preprint, Under Review, 2025

- We present Consistent Personalization for 3D Gaussian Splatting (CP-GS), a framework that progressively propagates the single-view reference appearance to novel perspectives, offering high-quality 3DGS personalization with faithful referential alignment.

Nautilus: Locality-aware Autoencoder for Scalable Mesh Generation
Yuxuan Wang*, Xuanyu Yi*, Haohan Weng*, Xiaokang Wei, Xianghui Yang, Chunchao Guo, Long Chen, Hanwang Zhang
International Conference on Computer Vision (ICCV), 2025

- We propose Nautilus, a locality-aware autoencoder for artist-like mesh generation, which leverages the local properties of manifold meshes to achieve structural fidelity and efficient representation.

View-Consistent 3D Editing with Gaussian Splatting
Yuxuan Wang, Xuanyu Yi, Zike Wu, Na Zhao, Long Chen, Hanwang Zhang
European Conference on Computer Vision (ECCV), 2024

- In the diffusion model, we proposed effective multi-view consistency designs that harmonize the inconsistent multi-view image guidance by integrating with 3D Gaussian Splatting (3DGS) characteristics, offering high-quality 3DGS editing.

Predicate Debiasing in Vision-Language Models Integration for Scene Graph Generation
Yuxuan Wang, Xiaoyuan Liu
Main Conference, Empirical Methods in Natural Language Processing (EMNLP), 2024

- We introduced a plug-and-play debiasing method for the zero-shot VLMs, dynamically ensembling them to address the underrepresentation issue in Scene Graph Generation (SGG) models.

GEB+: A Benchmark for Generic Event Boundary Captioning, Grounding and Retrieval
Yuxuan Wang, Difei Gao, Licheng Yu, Stan Weixian Lei, Matt Feiszli, Mike Zheng Shou
European Conference on Computer Vision (ECCV), 2022

- We introduced three tasks of video boundary understanding on our new dataset called Kinetics-GEB+ (Generic Event Boundary Plus), consisting of over 170k boundaries associated with captions in 12K videos.
- We designed a new Temporal-based Pairwise Difference (TPD) Modeling method for visual difference representation and achieved significant performance improvements.

Pushing Rendering Boundaries: Hard Gaussian Splatting
Qingshan Xu, Jiequan Cui, Xuanyu Yi, Yuxuan Wang, Yuan Zhou, Yew-Soon Ong, Hanwang Zhang
Arxiv Preprint, Under Review, 2025

- We propose Hard Gaussian Splatting, dubbed HGS, which considers multi-view significant positional gradients and rendering errors to grow hard Gaussians that fill the gaps of classical Gaussian Splatting on 3D scenes, thus achieving superior NVS results.

PBR3DGen: A VLM-guided mesh generation with high-quality PBR texture
Xiaokang Wei, Bowen Zhang, Xianghui Yang, Yuxuan Wang, Xi Zhao, Yan Luximon
Arxiv Preprint, Under Review, 2025

- We present PBR3DGen, a two-stage mesh generation method that produces high-quality PBR materials by integrating a novel multi-view PBR material estimation model and a PBR-based Large Reconstruction Model (PBR-LRM).

DragNeXt: Rethinking Drag-Based Image Editing
Yuan Zhou, Junbao Zhou, Qingshan Xu, Kesen Zhao, Yuxuan Wang, Hao Fei, Richang Hong, Hanwang Zhang
Arxiv Preprint, Under Review, 2025

- We propose a simple-yet-effective editing framework, dubbed DragNeXt, redefining Drag-Based Image Editing (DBIE) as deformation, rotation, and translation of user-specified handle regions.

Symbolic Replay: Scene Graph as Prompt for Continual Learning on VQA Task
Stan Weixian Lei, Difei Gao, Jay Zhangjie Wu, Yuxuan Wang, Wei Liu, Mengmi Zhang, Mike Zheng Shou
AAAI Conference on Artifical Intelligence (AAAI), 2023, **Oral**

- We introduced Scene Graph as Prompt (SGP) for symbolic replay, a real-data-free replay-based method for Continual Learning VQA, which overcomes the limitations of replay-based methods by leveraging the scene graph as an alternative to images for replay.

AssistSR: Task-oriented Video Segment Retrieval for Personal AI Assistant
Stan Weixian Lei, Difei Gao, Yuxuan Wang, Dongxing Mao, Zihan Liang, Lingmin Ran, Mike Zheng Shou
Findings, Empirical Methods in Natural Language Processing (EMNLP), 2022

- We introduce a new dataset and a new task called Affordance-centric Question-driven Video Segment Retrieval (AQVSR), aiming at retrieving affordance-centric instructional video segments given users' questions.
- To address the task, we developed a straightforward model called Dual Multimodal Encoders (DME).

INTERNSHIP EXPERIENCE

Research Intern | Tencent Hunyuan
Hunyuan 3D AIGC Center, Tencent TEG — Jun. 2024

Software Development Intern | Inspur
AI and Big Data Platform for Health Service — Jun. 2020

RESEARCH PROJECTS

IMS-MLD Decoding Algorithm for Reed-Muller Code
Senior Thesis supervised by Prof. *Qin Huang*, Beihang University — Feb. 2020

Blind Identification and Demodulation of Modulated Signals
Supervised by Prof. *Qin Huang*, Beihang University — Jul. 2018

SKILLS

- **Language**: Mandarin (Native), English (GRE: 331/340, TOEFL: 110/120)
- **Coding**: Python, Java, C, MATLAB, SQL
- **Music**: Piano, Guitar, Clarinet, Cavalry trumpet, Singing, Chorus Conducting